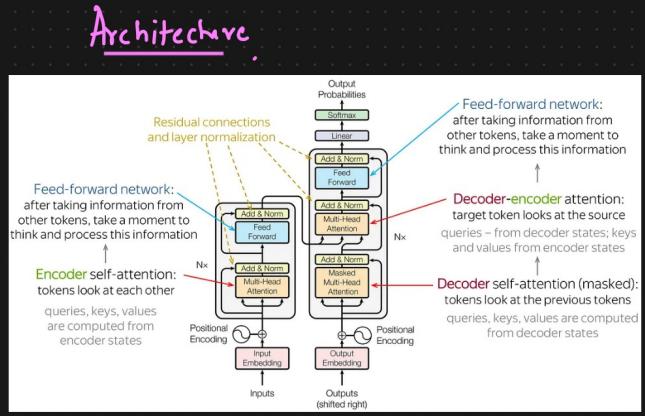


Introduction To Transformers

- 1) RNN / LSTM / GRU RNN
- 2) Encoder Decoder Architecture
- 3) ATTENTION MECHANISM
- 4) TRANSFORMERS

- 1) Why Transformers?
- 2) Architecture of Transformers?
- 3) SELF ATTENTION $\rightarrow Q, K, V$
- 4) Positional Encoding
- 5) Multi Head ATTENTION
- 6) Combining the Working of Transformers



Generative AI \rightarrow LM, Multimodal

BERT, GPT \leftarrow

Open AI \rightarrow ChatGPT

GPT-4o

① What And Why \rightarrow Transformers

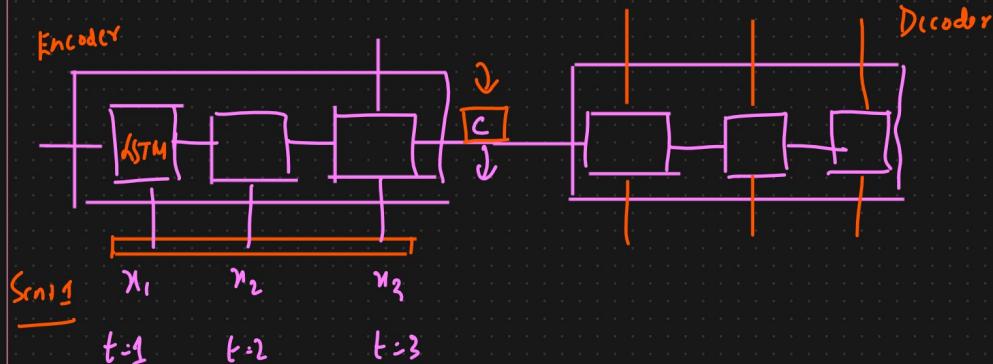
Transformers in natural language processing (NLP) are a type of deep learning model that use self-attention mechanisms to analyze and process natural language data. They are encoder-decoder models that can be used for many applications, including machine translation. \Rightarrow Seq2Seq Task

Eg: Language Translation \rightarrow Google Translation

English \rightarrow French

if \Rightarrow Many \rightarrow O/p: Many. $\{$ Length of the sentence $\}$.

Encoder - Decoder



Sentence Length $\uparrow\uparrow$

Beam Score $\downarrow\downarrow$

Length Sentence $\uparrow\uparrow$

3.1 DECODER: GENERAL DESCRIPTION

In a new model architecture, we define each conditional probability in Eq. (2) as:

$$p(y_i | y_1, \dots, y_{i-1}, \mathbf{x}) = g(y_{i-1}, s_i, c_i), \quad (4)$$

where s_i is an RNN hidden state for time i , computed by

$$s_i = f(s_{i-1}, y_{i-1}, c_i).$$

It should be noted that unlike the existing encoder-decoder approach (see Eq. (2)), here the probability is conditioned on a distinct context vector c_i for each target word y_i .

The context vector c_i depends on a sequence of annotations (h_1, \dots, h_{T_x}) to which an encoder maps the input sentence. Each annotation h_i contains information about the whole input sequence with a strong focus on the parts surrounding the i -th word of the input sequence. We explain in detail how the annotations are computed in the next section.

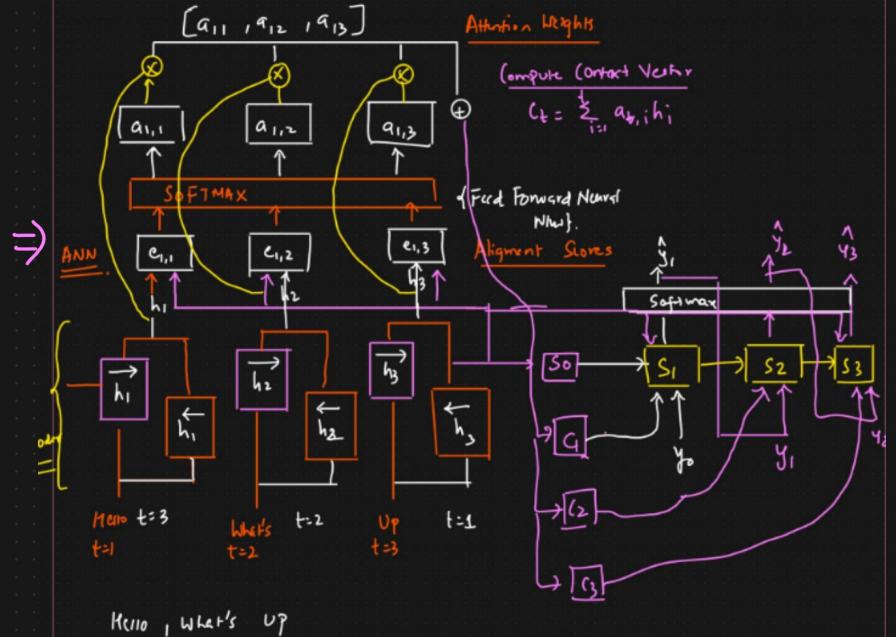
The context vector c_i is, then, computed as a weighted sum of these annotations h_i :

$$c_i = \sum_{j=1}^{T_x} \alpha_{ij} h_j. \quad (5)$$

Figure 1: The graphical illustration of the proposed model trying to generate the i -th target word y_i given a source sentence (x_1, x_2, \dots, x_T) .

Additional Context \rightarrow Decoder

long Sentence. Accuracy \uparrow .



Attention Mechanism

① Parallelly we cannot send all the words in a sentence \rightarrow Scalable

DATASET \rightarrow Huge \rightarrow Scalable with Respect to Training.

TRANSFORMERS \neq LSTM RNN

Self Attention Module \leftarrow All the words will be parallelly sent to encoder.

\Downarrow

Positional Encoding

Transformer \uparrow DATASET \rightarrow Amazing SOTA \leftarrow NLP

Transfer Learning \rightarrow MultiModal Task \rightarrow NLP + Image \leftarrow

Transformers \div AI Space \rightarrow SOTA Model \rightarrow

$\downarrow \checkmark \downarrow \downarrow$

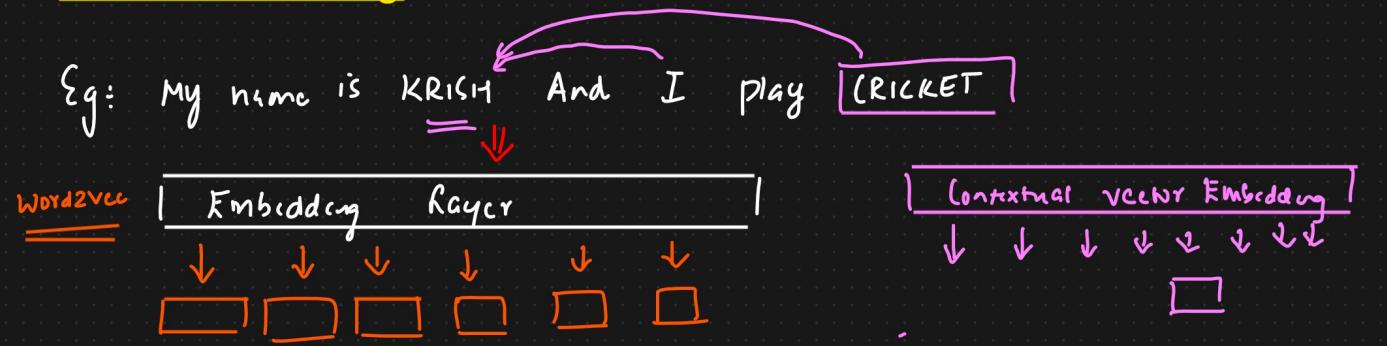
BERT GPT \rightarrow Transfer learning \rightarrow SOTA Models \rightarrow DALLE \leftarrow Generating AI

Train huge Data

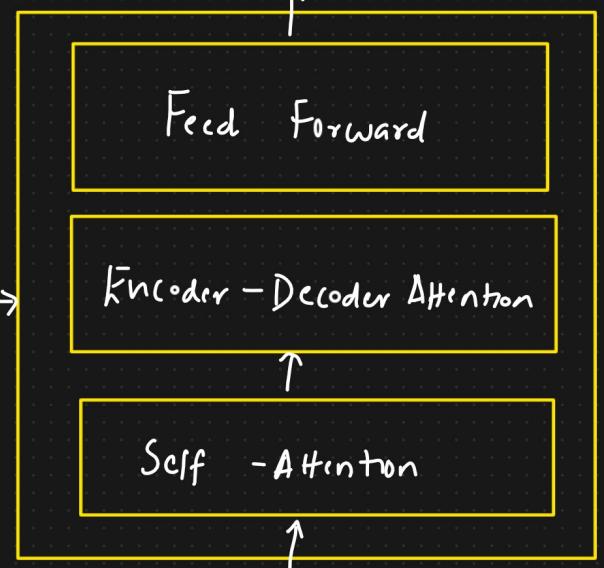
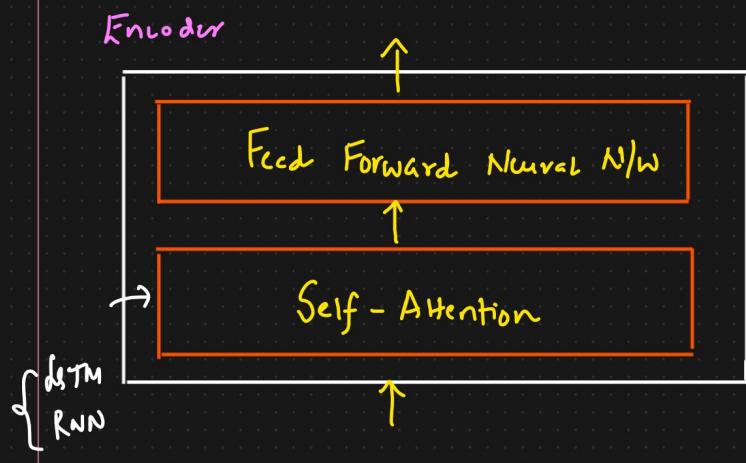
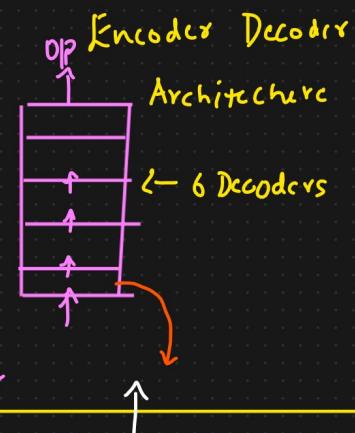
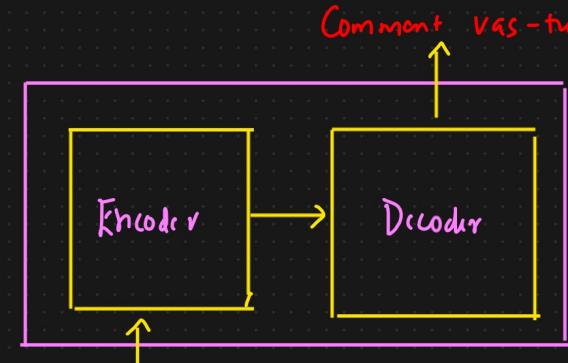
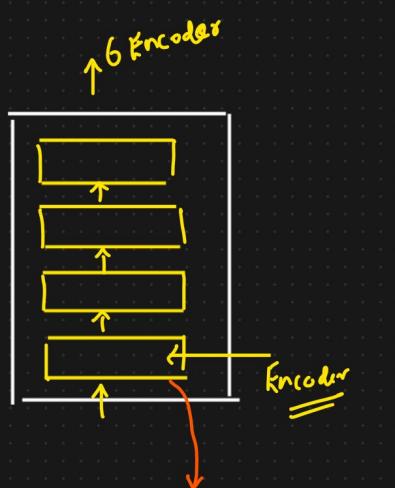
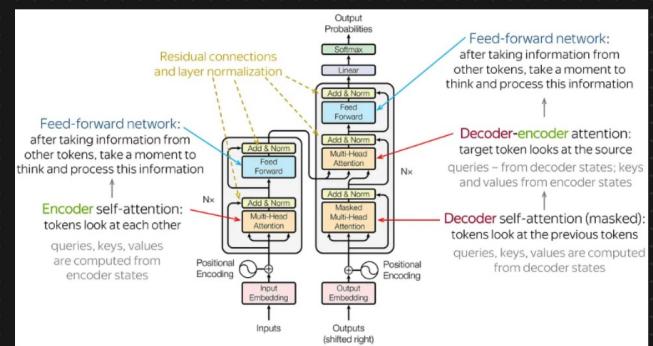
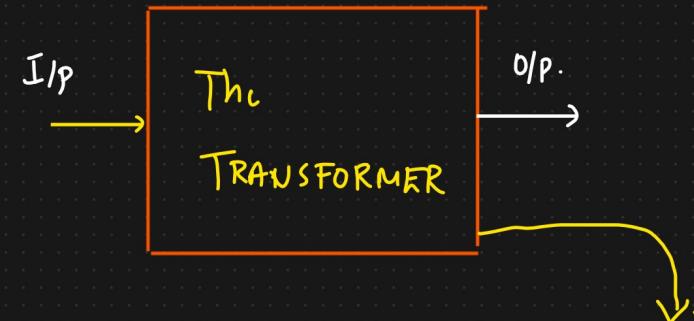
AI's

Generating AI

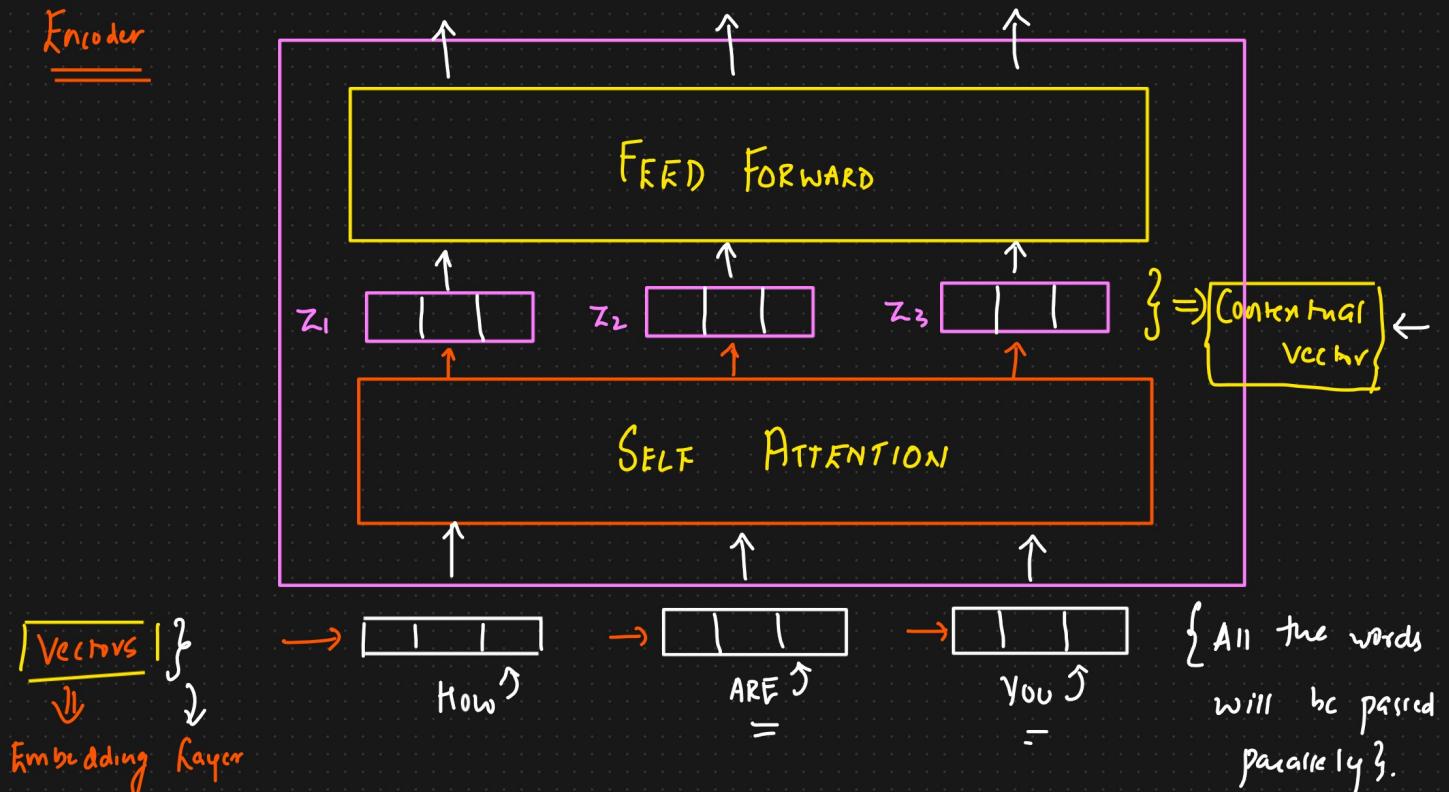
② Contextual Embedding → Self Attention



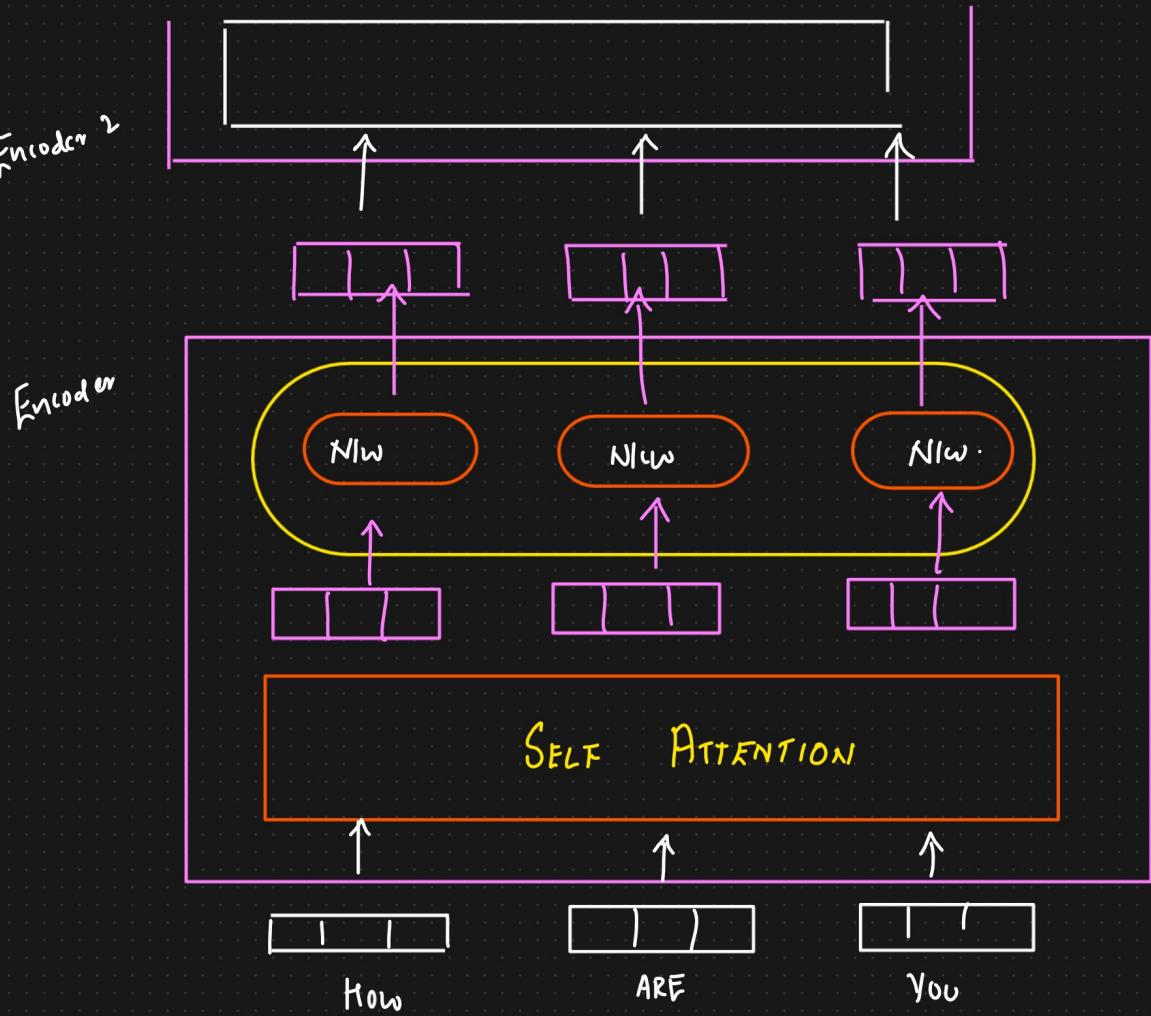
② Basic Transformer Architecture {Seq2Seq Task} → Language Translation {Eng → French}



Encoder



Encoder 2



Self Attention At a Higher Level

Eg: The cat sat on the mat, the cat lay on the rug.

Word Embedding
↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓

SELF ATTENTION

↓
↓↓
↓

The. The

→ [Cat] → [Cat] → [] [] []
1 → [Sat] Sat

{ Contextual Embedding }

Rank 2 → on
= 3 → the
mat

mat → [- | - | -]



Self Attention In Detail

① To Create 3 vectors from each of the encoder i/p. Query vector,

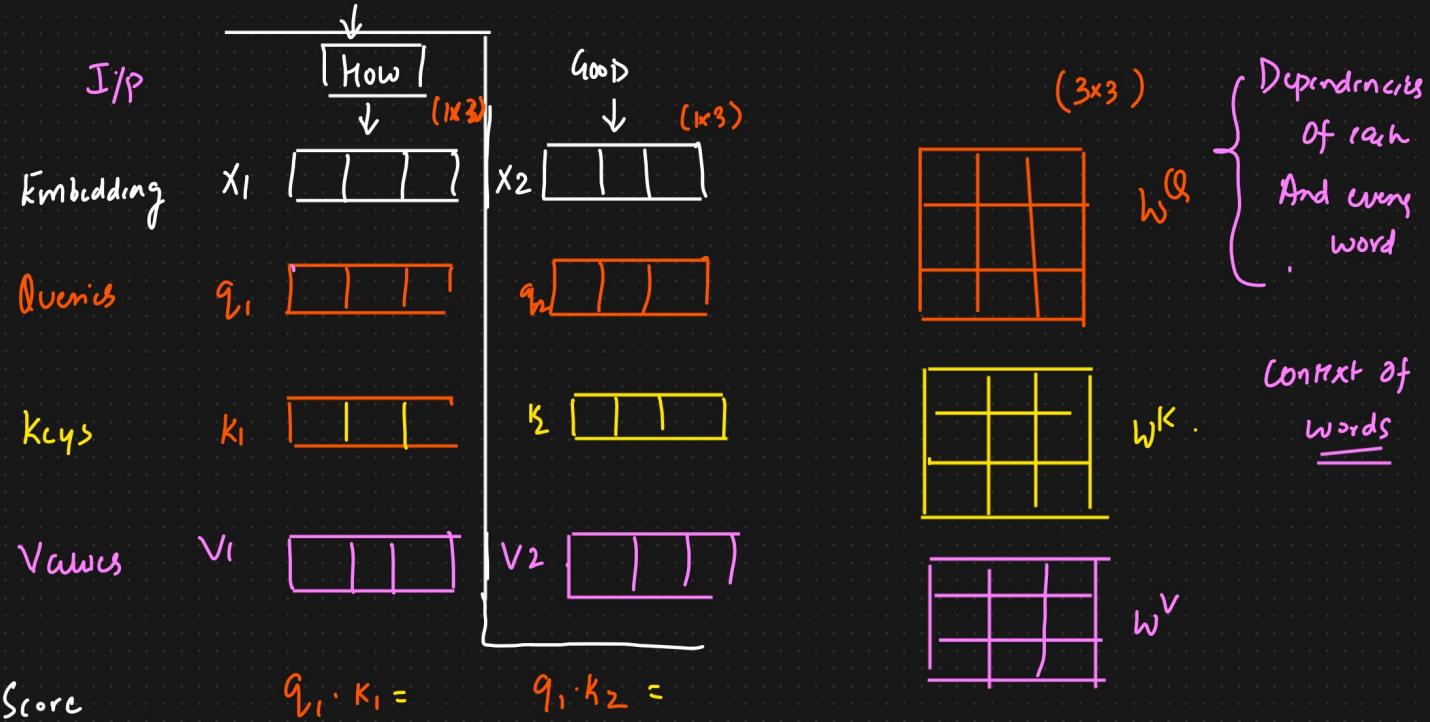
Key vector, value vector. \Rightarrow Contextual Embedding

Eg: YT \Rightarrow Search keywords \Rightarrow {Query}.

Query \rightarrow {Key} \rightarrow Tags
 \rightarrow Description \Rightarrow Value \Rightarrow o/p Video
 \Rightarrow Title

② Second step in calculating self attention is to calculate the Score.

The Score determines how much focus to place on the other part of the sentence.



$$\text{Score} \quad q_1 \cdot k_1 = \quad q_1 \cdot k_2 =$$

④ How much focus to place on other part of the I/P Sentence as we encode a word at certain position

Divide \sqrt{dk} — — Dimension $3 \uparrow$ Dimension $12 \uparrow$

More stable \downarrow O/P $\uparrow \uparrow \uparrow$

gradients $\frac{\text{Variance}}{\sqrt{dk}} \approx \frac{\text{Variance}}{\sqrt{dk}} \Rightarrow$

SoftMax $\boxed{[0-1]} \quad \boxed{[0-1]} \cdot \boxed{[0-1]}$

$0.88 + 0.12 = 1$

Softmax \Rightarrow The Softmax score determines how much each word will be expressed at this position

X X X

Value Vectors $v_1 \boxed{\quad \quad}$ $v_2 \cdot \boxed{1 \quad 0.88 \quad 0.12}$

$0.88 \boxed{1 \quad 1 \quad 1}$

$\frac{0.88 \boxed{0.88 \quad 0.88}}{0.88 + 0.12} + \frac{0.12 \boxed{0.6 \quad 0.12}}{0.88 + 0.12}$

Contextual Vector $\xrightarrow{I_1} \boxed{1 \quad 1 \quad 1}$ $\xrightarrow{I_2} \boxed{1 \quad 1 \quad 1}$

Self Attention At Higher And Detailed Level

Self-attention, also known as scaled dot-product attention, is a crucial mechanism in the transformer architecture that allows the model to weigh the importance of different tokens in the input sequence relative to each other.

Idea :

→ $[0.8 | 0.3 | 0.4]$ $[0.4 | 0.6 | 0.2]$ $[0.3 | 0.4 | 0.5]$ { Contextual Embedding }

Language Translation $\boxed{\text{Self Attention}}$

Text Summarization \uparrow \uparrow \uparrow { Sentence, Dataset }.

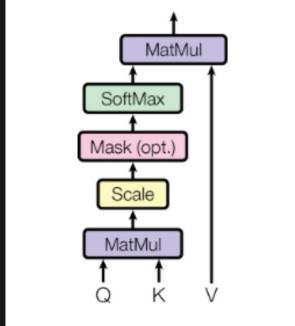
$[1 | 1 | 0 | 0]$ $[0 | 1 | 1 | 0]$ $[0 | 0 | 0 | 1]$

The CAT SAT

Embedding

layer → Fixed Vector

Scaled Dot-Product Attention



1) Inputs: Queries, Keys, And Values

Model → Queries, Keys And Values

1. Query Vectors (Q):

Role: Query vectors represent the token for which we are calculating the attention. They help determine the importance of other tokens in the context of the current token.

Importance:

Focus Determination: Queries help the model decide which parts of the sequence to focus on for each specific token. By calculating the dot product between a query vector and all key vectors, the model assesses how much attention to give to each token relative to the current token.

Contextual Understanding: Queries contribute to understanding the relationship between the current token and the rest of the sequence, which is essential for capturing dependencies and context.

2. Key Vectors (K):

Role: Key vectors represent all the tokens in the sequence and are used to compare with the query vectors to calculate attention scores.

Importance:

Relevance Measurement: Keys are compared with queries to measure the relevance or compatibility of each token with the current token. This comparison helps in determining how much attention each token should receive.

Information Retrieval: Keys play a critical role in retrieving the most relevant information from the sequence by providing a basis for the attention mechanism to compute similarity scores.

3. Value Vectors (V):

Role: Value vectors hold the actual information that will be aggregated to form the output of the attention mechanism.

Importance:

Information Aggregation: Values contain the data that will be weighted by the attention scores. The weighted sum of values forms the output of the self-attention mechanism, which is then passed on to the next layers in the network.

Context Preservation: By weighting the values according to the attention scores, the model preserves and aggregates relevant context from the entire sequence, which is crucial for tasks like translation, summarization, and more.

$$\text{Input Sequence} = \left[\text{"The"}, \text{"CAT"}, \text{"SAT"} \right]$$

$$\text{Embedding Size} = 4$$

$$Q, K, V \Rightarrow 4$$

$$\xrightarrow{\text{L}} \boxed{\text{S} \text{ H} \text{ A} \text{ T}} \rightarrow \boxed{\text{O} \text{ P} \text{ I} \text{ O} \text{ S}} \leftarrow$$

$$\boxed{\text{I} \text{ P} \text{ O}} \quad \boxed{\square} \quad \boxed{\square}$$

① Token Embedding

$$E_{\text{The}} = [1 \ 0 \ 1 \ 0]$$

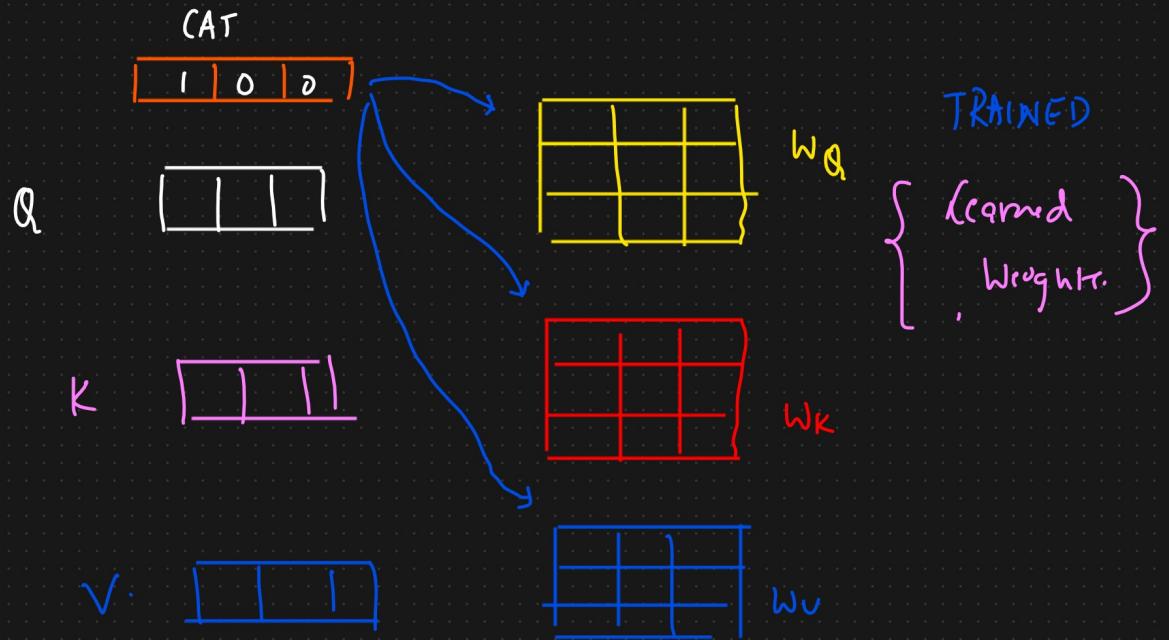
$$E_{\text{SAT}} = [1, 1, 1, 1]$$

$$E_{\text{CAT}} = [0 \ 1 \ 0 \ 1]$$

Sentence, Dataset

② Linear Transformation

We create Q, K, V by multiplying the embeddings by learned weights matrices W_Q , W_K and W_V .



lets consider

$$W_Q = W_K = W_V = I \quad \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \leftarrow$$

$$Q_{\text{The}} = [1 0 1] \cdot \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} = [1 0 1]$$

$$K_{\text{The}} = [1 0 1]$$

$$V_{\text{The}} = [1 0 1]$$

$$\textcircled{1} \quad Q_{\text{The}} = K_{\text{The}} = V_{\text{The}} = [1 0 1]$$

$$\textcircled{2} \quad Q_{\text{CAT}} = K_{\text{CAT}} = V_{\text{CAT}} = [0 1 0 1]$$

$$\textcircled{3} \quad Q_{\text{SAT}} = K_{\text{SAT}} = V_{\text{SAT}} = [1, 1, 1, 1]$$

③ Compute Attention Scores

$$\begin{bmatrix} 1 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} 1 \\ 0 \\ 1 \\ 0 \end{bmatrix}$$

The

$$\text{Score}(Q_{\text{The}}, K_{\text{The}}) = \begin{bmatrix} 1 & 0 & 1 & 0 \end{bmatrix} \cdot \begin{bmatrix} 1 & 0 & 1 & 0 \end{bmatrix}^T = 2$$

$$\text{Score}(Q_{\text{The}}, K_{\text{CAT}}) = \begin{bmatrix} 1 & 0 & 1 & 0 \end{bmatrix} \cdot \begin{bmatrix} 0 & 1 & 0 & 1 \end{bmatrix}^T = 0$$

$$\text{Score}(Q_{\text{The}}, K_{\text{SAT}}) = \begin{bmatrix} 1 & 0 & 1 & 0 \end{bmatrix} \cdot \begin{bmatrix} 1 & 1 & 1 & 1 \end{bmatrix}^T = 2$$

For the token CAT

$$\text{Score}(Q_{\text{CAT}}, K_{\text{The}}) = \begin{bmatrix} 0 & 1 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} 1 & 0 & 1 & 0 \end{bmatrix}^T = 0$$

$$\text{Score}(Q_{\text{CAT}}, K_{\text{CAT}}) = \begin{bmatrix} 0 & 1 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} 0 & 1 & 0 & 1 \end{bmatrix}^T = 2$$

$$\text{Score}(Q_{\text{CAT}}, K_{\text{SAT}}) = \begin{bmatrix} 0 & 1 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} 1 & 1 & 1 & 1 \end{bmatrix}^T = 2.$$

For the Token SAT

$$\text{Score}(Q_{\text{SAT}}, K_{\text{The}}) = \begin{bmatrix} 1 & 1 & 1 & 1 \end{bmatrix} \cdot \begin{bmatrix} 1 & 0 & 1 & 0 \end{bmatrix}^T = 2$$

$$\text{Score}(Q_{\text{SAT}}, K_{\text{CAT}}) = 2$$

$$\text{Score}(Q_{\text{SAT}}, K_{\text{SAT}}) = 4$$

④ Scaling

We take up the scores and scale down by dividing the scores by the

$$\sqrt{d_K} \Rightarrow d_K = 4 \quad \sqrt{d_K} = 2.$$

Scaling in the attention mechanism is crucial to prevent the dot product from growing too large. \Rightarrow Ensure stable gradients during Training.

d_K is large \rightarrow

① Gradient Exploding

② Softmax Saturation $\{\curvearrowright\}$ \rightarrow Vanishing Gradient Problem.

$$Q = \begin{bmatrix} 2 & 3 & 4 & 1 \end{bmatrix} \quad K_1 = \begin{bmatrix} 1 & 0 & 1 & 0 \end{bmatrix} \quad K_2 = \begin{bmatrix} 0 & 1 & 0 & 1 \end{bmatrix}$$

Without Scaling

$$Q \cdot K_1^T = 2 \times 1 + 3 \times 0 + 4 \times 1 + 1 \times 0 = 2 + 4 = 6.$$

$$Q \cdot K_2^T = 2 \times 0 + 3 \times 1 + 4 \times 0 + 1 \times 1 = 0 + 3 + 0 + 1 = 4$$

* Score $[6, 4] \Rightarrow$ Scaling Not Applied

$$\text{Softmax}([6, 4]) = \left[\frac{e^6}{e^6 + e^4}, \frac{e^4}{e^6 + e^4} \right] = \left[\frac{e^6}{e^6(1 + e^{-2})}, \frac{e^4}{e^4(e^2 + 1)} \right]$$

① Property of Softmax WQ, WK, WV

$$([10, 1]) = \left[\frac{0.99}{\cancel{10}}, \frac{0.01}{\cancel{1}} \right]$$

$$= \left[\frac{1}{(1 + e^{-2})}, \frac{1}{(e^2 + 1)} \right]$$

$$\text{Dot product} = \text{Large values} \approx [0.88, 0.12].$$

Most of the attention weight is assigned to the first key vector, very little to the second vector,

With Scaling

① Compute Scaled Dot Product

$$[6, 4] \Rightarrow \text{Scale} \Rightarrow \left[\frac{6}{\sqrt{2}}, \frac{4}{\sqrt{2}} \right] = \left[3, 2 \right]. \quad \frac{\sqrt{dk}}{\sqrt{dk}} \uparrow \text{same}$$

$$\sqrt{dk} = \sqrt{4} = \sqrt{2} \Rightarrow \text{Variance } 2, 3$$

dimension \uparrow Variance $\uparrow 2$

$$\text{Softmax}([3, 2]) = \left[\frac{e^3}{e^3 + e^2}, \frac{e^2}{e^3 + e^2} \right] = \left[\frac{e^3}{e^3(1 + e^{-1})}, \frac{e^2}{e^2(e^1 + 1)} \right]$$

$$= [0.73, 0.27] \Rightarrow \text{Attention weights}$$

④ Here, the attention weights are more balanced compared to the unscaled case

Summary of Importance

Stabilizing Training: Scaling prevents extremely large dot products, which helps in stabilizing the gradients during backpropagation, making the training process more stable and efficient.

Preventing Saturation: By scaling the dot products, the softmax function produces more balanced attention weights, preventing the model from focusing too heavily on a single token and ignoring others.

Improved Learning: Balanced attention weights enable the model to learn better representations by considering multiple relevant tokens in the sequence, leading to better performance on tasks that require context understanding.

Scaling ensures that the dot products are kept within a range that allows the softmax function to operate effectively, providing a more balanced distribution of attention weights and improving the overall learning process of the model.

$$④ \text{Scaling} = \sqrt{d_K} = \sqrt{4} \Rightarrow 2$$

Similarly Scaling

$$\text{Scaled-Score} (Q_{\text{The}}, K_{\text{The}}) = 2/2 = 1$$

will be done for
all other tokens.

$$\text{Scaled-Score} (Q_{\text{The}}, K_{\text{CAT}}) = 0/2 = 0$$

$$\text{Scaled-Score} (Q_{\text{The}}, K_{\text{SAT}}) = 2/2 = 1$$

⑤ Apply Softmax

$$\text{ATTENTION WEIGHTS}_{\text{"The"}} = \text{Softmax} ([1, 0, 1]) = [0.4223, 0.1554, 0.4223]$$

$$\text{ATTENTION WEIGHTS}_{\text{"CAT"}} = \text{Softmax} ([0, 2, 2]) = [0.1554, 0.4223, 0.4223]$$

$$\text{ATTENTION WEIGHTS}_{\text{"SAT"}} = \text{Softmax} ([2, 2, 4]) = [0.2119, 0.2119, 0.5762]$$

⑥ Weight Sum of Values

We multiply the attention weights by corresponding value vectors

For the Token The =

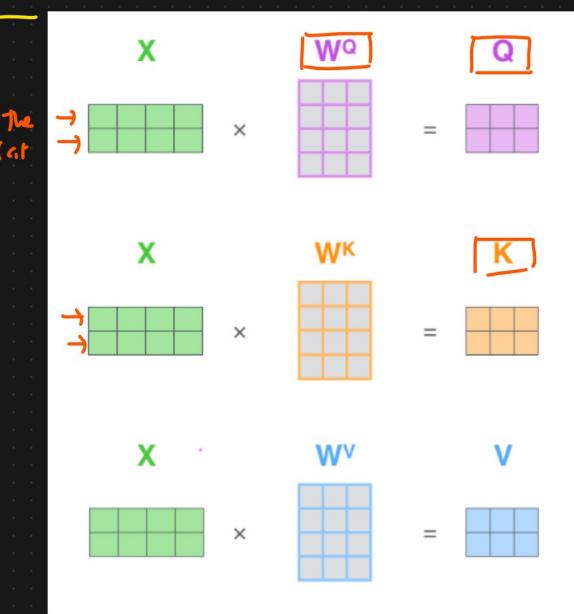
$$\begin{aligned}\text{Output}_{(\text{The})} &= 0.4223 * V_{\text{The}} + 0.1554 * V_{\text{CAT}} + 0.4223 * V_{\text{Sal}} \\ &= 0.4223 [1 0 1 0] + 0.1554 [0 1 0 1] + 0.4223 [1 1 1] \\ &= [0.4223, 0, 0.4223, 0] + [0, 0.1554, 0, 0.1554] + [0.4223, 0.4223, \\ &\quad 0.4223, 0.4223] \\ &= [1.2669, 0.9999, 1.2669, 0.9999].\end{aligned}$$

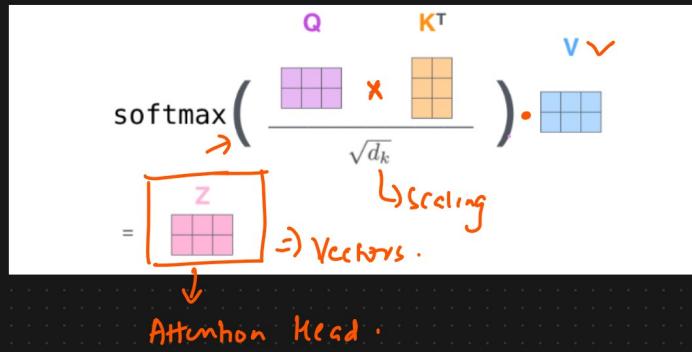
↓ Contextual
vector

The 1 1 0 1 1 0 \Rightarrow Self Attention $\Rightarrow [1.2669, 0.9999, 1.2669, 0.9999]$.

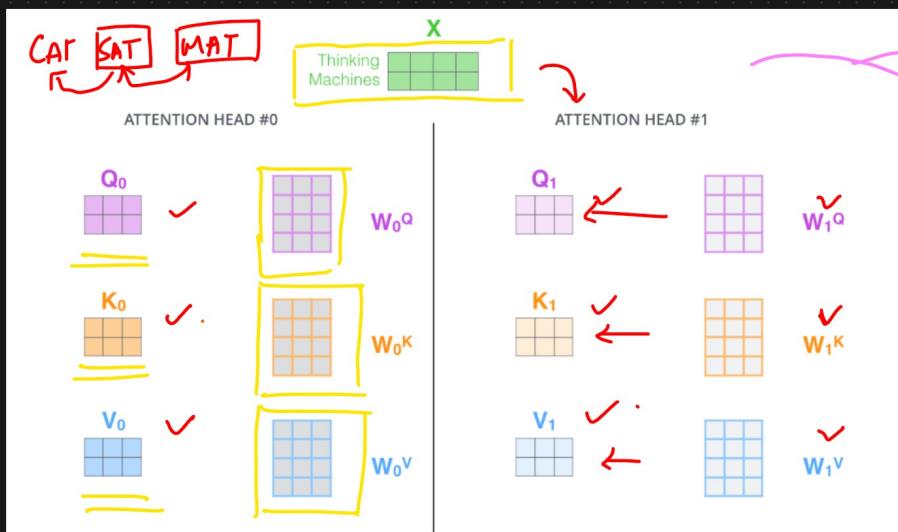
① $\hookrightarrow Q, K, V [w^Q, w^K, w^V]$
② \hookrightarrow Attention Score
③ \hookrightarrow Scaled
④ \hookrightarrow Softmax
⑤ \hookrightarrow Weighted Sum of Value (Softmax \times V)

④ Multi Head Attention





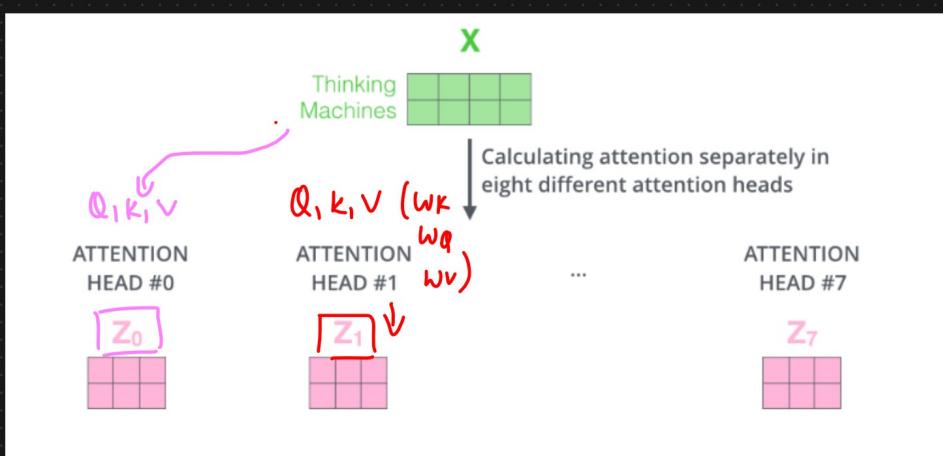
→ Self Attention with Multi Heads



It expands model's ability to focus on different position of tokens.

Score, Softmax $\times V$
 \downarrow \downarrow \downarrow \downarrow
 $[\dots]$ $[\dots]$ $[\dots]$ $[\dots]$
 Z_0 Z_1 \dots
 $\underline{\text{Vectors}}$

Multi Head Attention

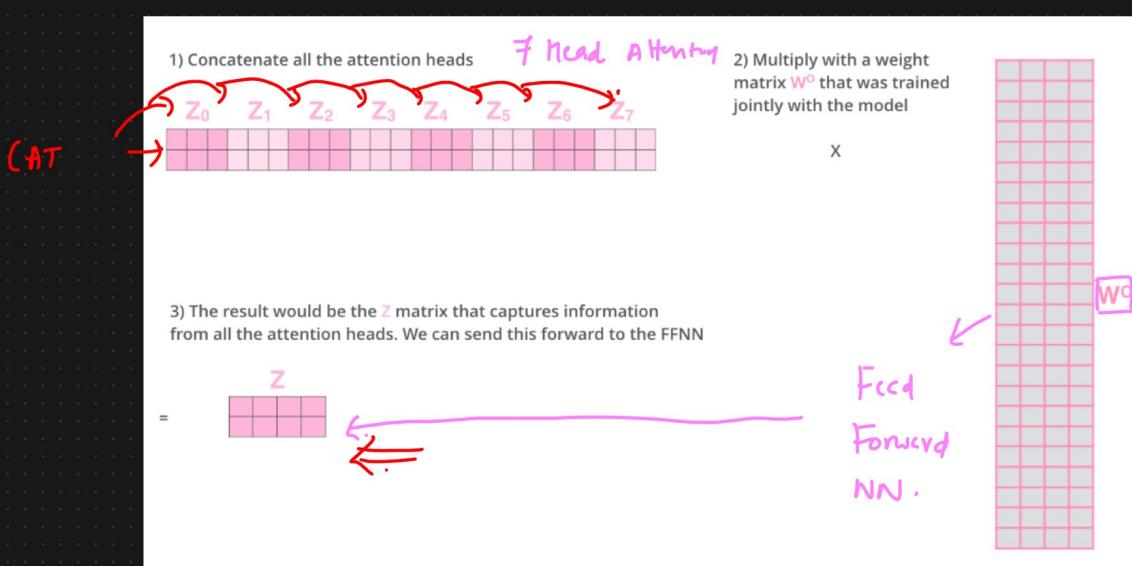
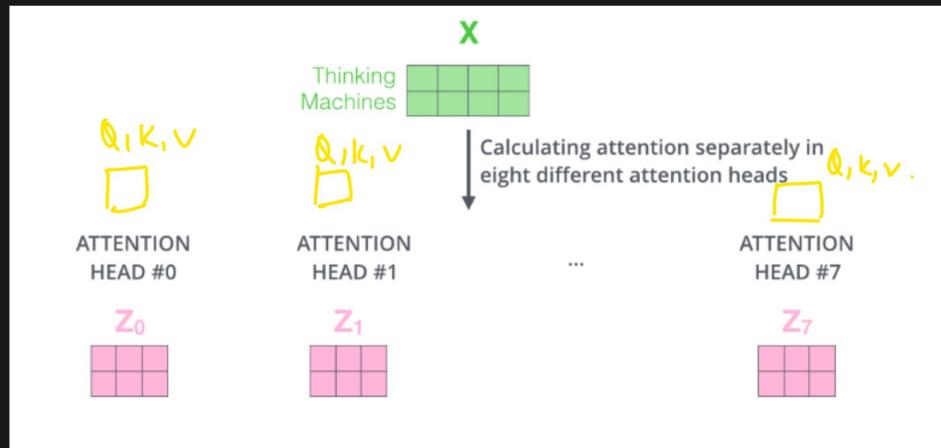
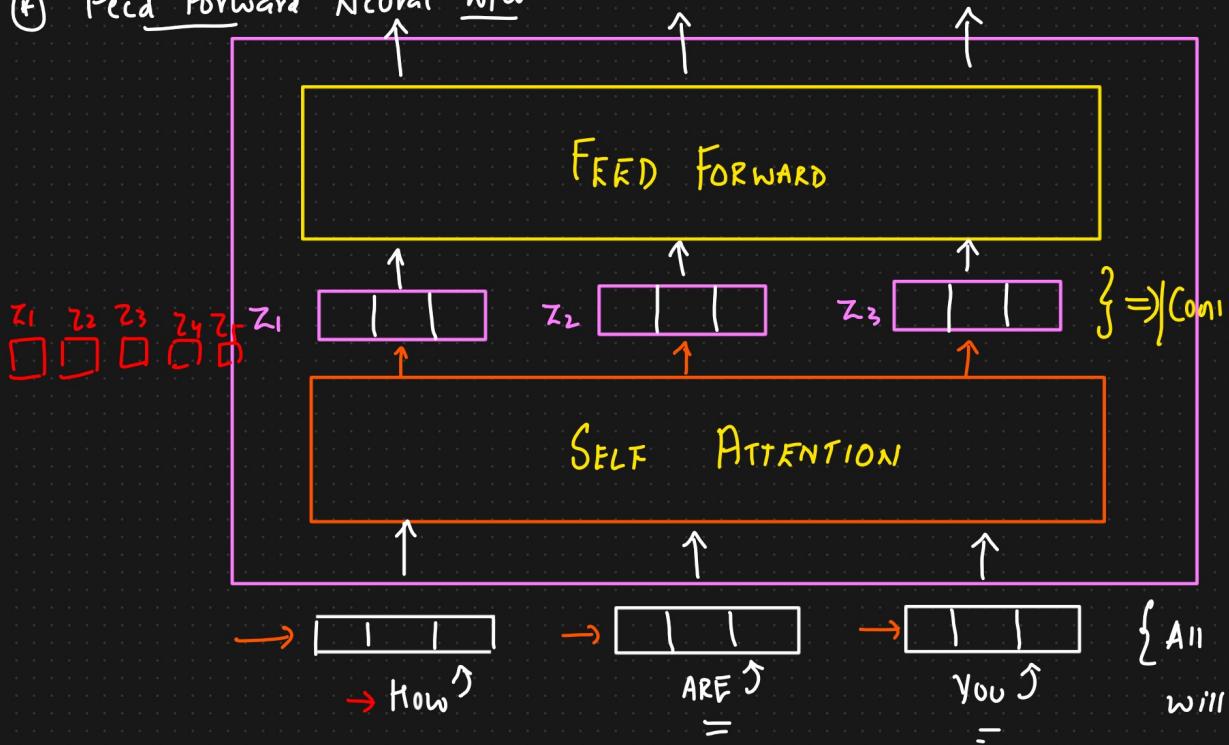


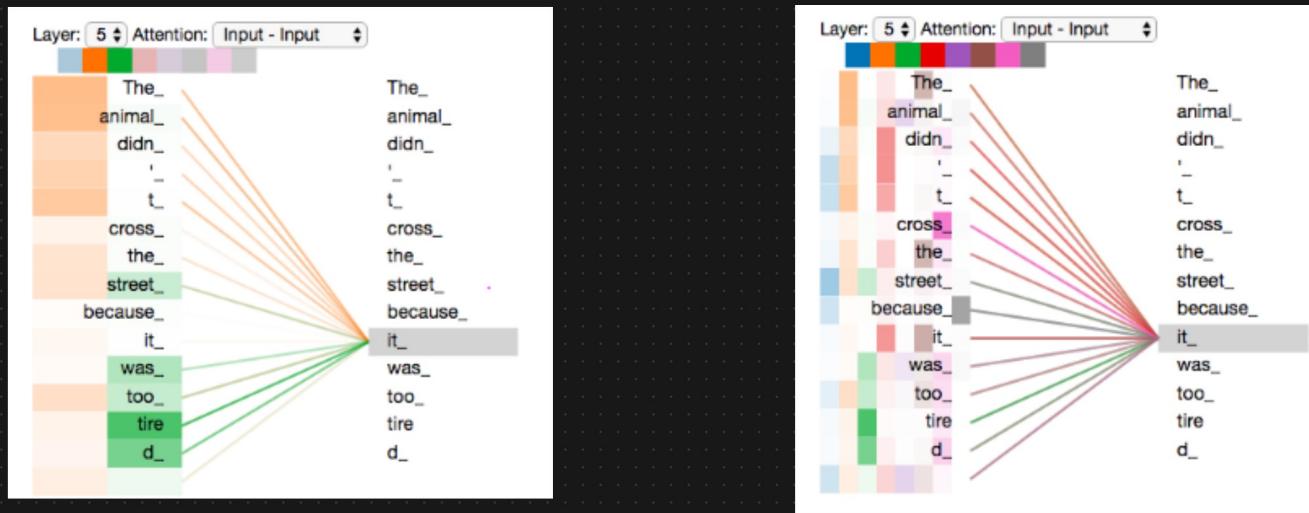
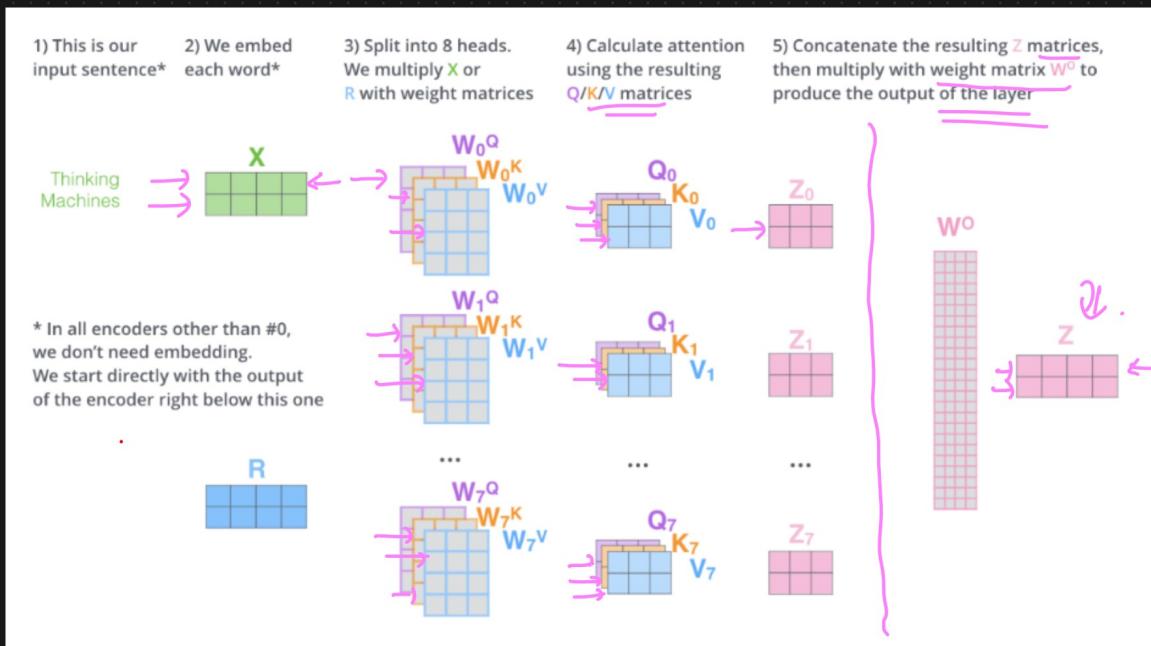
$[Z]$

$[Z]$

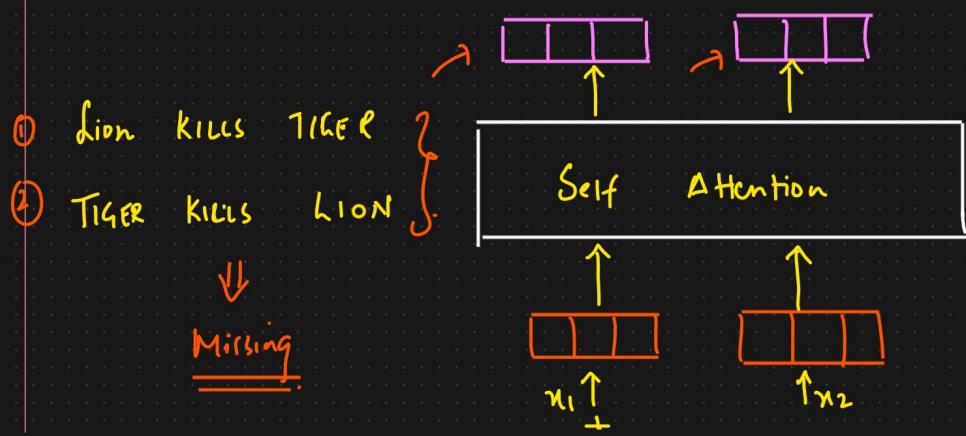
$[Z]$

⑥ Feed Forward Neural Network





⑥ Positional Encoding - Representing Order of Sequence



Advantage

- ① Word Tokens it can process parallelly
- ↓

DRAW BACK

lack the sequential structure of the words
{order}

Attention IS All

YOU NEED



Positional

Encoded

Vector

Book, Journal

Novel

↳ 1 lakh
words



Back propagation



Types of Position Encoding

1) Sinusoidal Position Encoding →

2) Learned Positional Encoding \Rightarrow Positional Encoding are learned during Training ↵

① Sinusoidal Positional Encoding \div It uses Sinc and Cosine functions

of different frequencies to create positional encodings

Formula \div



$$P.E(pos, 2i) = \sin\left(\frac{pos}{10000^{2i/d_{model}}}\right) \quad \text{where } pos \text{ is the position}$$

i is the dimension

$$P.E(pos, 2i+1) = \cos\left(\frac{pos}{10000^{2i/d_{model}}}\right) \quad d_{model} \text{ is the dimensionality}$$

of the embeddings.

Eg: The Cat Sat

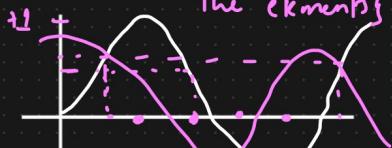
$$\text{The} \rightarrow [0.1 \ 0.2 \ 0.3 \ 0.4]$$

$$\text{CAT} \rightarrow [0.5 \ 0.6 \ 0.7 \ 0.8]$$

$$\text{SAT} \rightarrow [0.9 \ 1.0 \ 1.1 \ 1.2]$$



{ Miss the order of
the elements }



$$PE(pos, 2i) = \sin\left(\frac{pos}{10000^{2i/d_{model}}}\right) \quad PE(pos, 2i+1) = \cos\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

For our example $d_{model} = 4$

For position $pos = 0$

$$PE(0,0) = \sin\left(\frac{0}{10000^{0/4}}\right) = \sin(0) = 0$$

$$PE(0,1) = \cos\left(\frac{0}{10000^{1/4}}\right) = \cos(0) = 1$$

$$PE(0,2) = \sin\left(\frac{0}{10000^{2/4}}\right) = \sin(0) = 0$$

$$PE(0,3) = \cos\left(\frac{0}{10000^{3/4}}\right) = 1$$

$$PE = [0, 1, 0, 1]$$

$$PE(pos, 2i) = \sin\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

For $pos = 1$

$$PE(1,0) = \sin\left(\frac{1}{10000^{0/4}}\right) = \sin(1) = 0.8415$$

$$PE(1,1) = \cos\left(\frac{1}{10000^{1/4}}\right) \approx 0.5403$$

$$PE(1,2) = \sin\left(\frac{1}{10000^{2/4}}\right) \approx 0.01$$

$$PE(1,3) = \cos\left(\frac{1}{10000^{3/4}}\right) \approx 0.99995$$

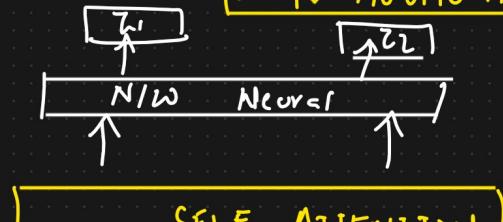
$$PE(pos, 2i+1) = \cos\left(\frac{pos}{10000^{2i+1/d_{model}}}\right)$$

Vector

$$\text{The } [0.1 \mid 0.2 \mid 0.3 \mid 0.4] \longrightarrow [0 \mid 1 \mid 0 \mid 1 \mid]$$

Positional Encoding

$$\text{CAT } [0.5 \mid 0.8 \mid 0.7 \mid 0.8] \longrightarrow [0.84 \mid 0.54 \mid 0.01 \mid 0.995]$$



→ Multi-head Attention

Diagram illustrating the addition of positional encodings to word embeddings:

Positional Encoding (PE) for 'CAT': $\begin{bmatrix} 0.84 & 0.34 & 0.01 & 0.99 \end{bmatrix}$

Positional Encoding (PE) for 'the': $\begin{bmatrix} 0 & 1 & 0 & 1 \end{bmatrix}$

Word Embedding (WE) for 'CAT': $\begin{bmatrix} 0.5 & 0.6 & 0.7 & 0.8 \end{bmatrix}$

Word Embedding (WE) for 'the': $\begin{bmatrix} 0.1 & 0.2 & 0.3 & 0.4 \end{bmatrix}$

The final embeddings are calculated as:

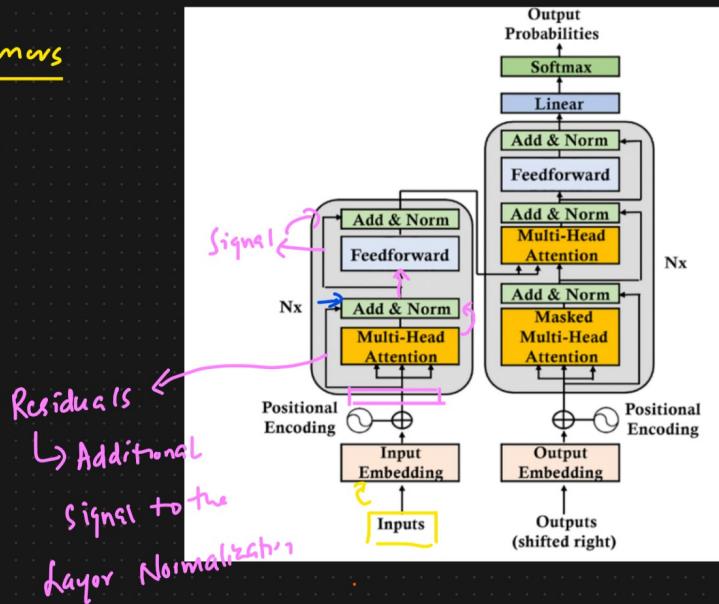
For 'CAT': $\begin{bmatrix} 0.5 & 0.6 & 0.7 & 0.8 \end{bmatrix} + \begin{bmatrix} 0.84 & 0.34 & 0.01 & 0.99 \end{bmatrix} = \begin{bmatrix} 1.34 & 0.94 & 0.01 & 1.79 \end{bmatrix}$

For 'the': $\begin{bmatrix} 0.1 & 0.2 & 0.3 & 0.4 \end{bmatrix} + \begin{bmatrix} 0 & 1 & 0 & 1 \end{bmatrix} = \begin{bmatrix} 0.1 & 1.2 & 0.3 & 0.4 \end{bmatrix}$

④ Layer Normalization In Transformers

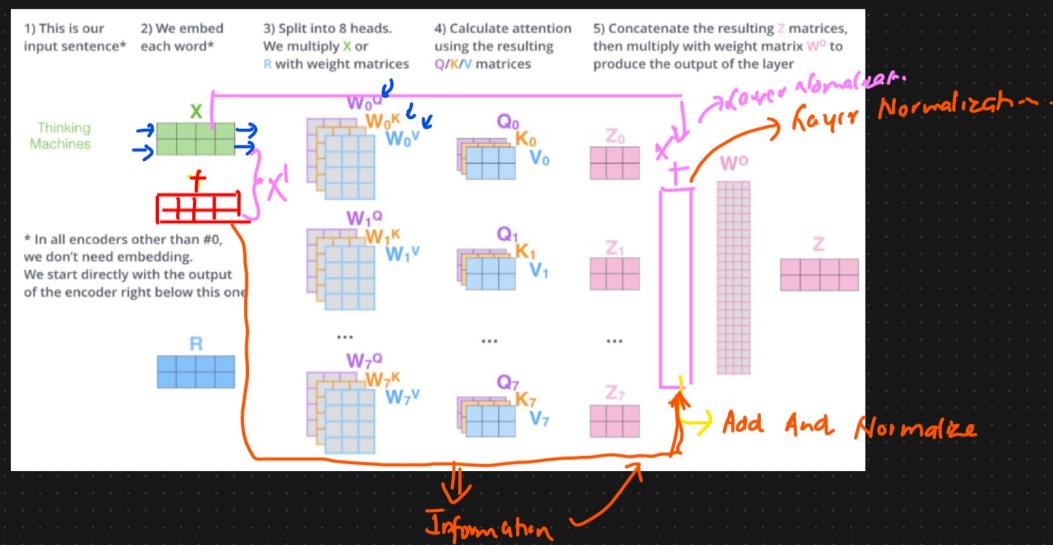
Transformers

Residuals



- ① Self Attention layer
 - ② Multi head Attention
 - ③ Positional Encoding
 - ④ Layer Normalization

ADD AND Normalize



Normalization

→ Batch Normalization

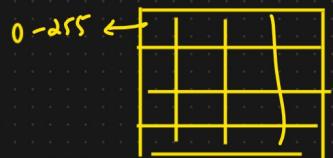
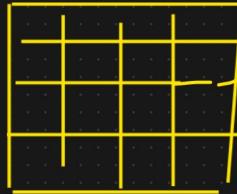
→ Layer Normalization



<u>f1</u>	<u>f2</u>	<u>House Size</u>	<u>No. of Rooms</u>	<u>Price</u>
1200	2			45
1500	3			70

Normalization : Standard Scaling

$$z_{score} = \frac{x_i - \mu}{\sigma} \Rightarrow \{\mu=0, \sigma=1\} \quad f_1 \rightarrow f_1' \Leftarrow \mu=0, \sigma=1$$

Deep learning : I/P Image \Rightarrow Min Max Scaler \Rightarrow Min Max Scaler \Rightarrow Advantages

① Improved Training Stability

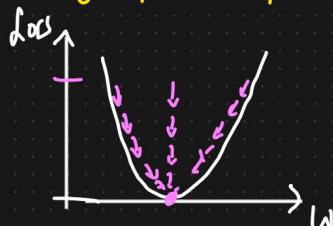
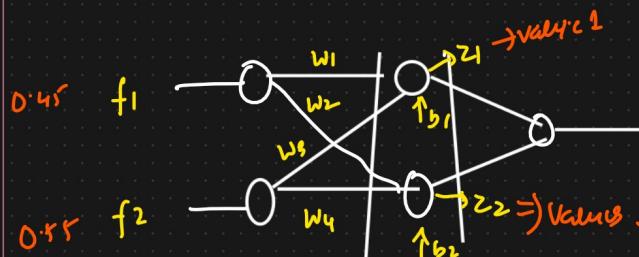


Vanishing and exploding Gradient problem

$$\left\{ \begin{array}{l} \mu=0 \\ \sigma=1 \end{array} \right.$$



② Faster Convergence

 \Rightarrow Back propagation \Rightarrow Stable update. f_1

$$\begin{bmatrix} \text{Mouse size} \\ \text{Rooms} \end{bmatrix} \rightarrow \begin{bmatrix} 0.45 \\ 0.60 \\ - \endbmatrix$$

 f_2

$$\begin{bmatrix} \text{Price} \\ z_1 \\ z_2 \end{bmatrix} \rightarrow \begin{bmatrix} 45 \\ - \\ - \\ - \end{bmatrix}$$

Normalization (Standard)

$$z_1 = \sigma \left[(0.45 \cdot w_1 + 0.55 \cdot w_3) + b_1 \right] = \text{Value 1} \quad \text{Does this distribution}$$

$$z_2 = \sigma \left(\dots + b_2 \right) = \text{Value 2} \quad \left. \begin{array}{l} \mu=0, \sigma=1 \\ \mu_1, \sigma_1 \\ \mu_2, \sigma_2 \end{array} \right\}$$

Batch Normalization VS Layer Normalization f_1 f_2 $z_1 \quad z_2$ $-$ $-$ $- \sqrt{1}$

$$z_{score} = \frac{x_i - \mu_1}{\sigma_1} \quad \left. \begin{array}{l} \mu_1, \sigma_1 \\ \mu_2, \sigma_2 \\ \mu_3, \sigma_3 \end{array} \right\}$$

 $-$ $-$ $-$

$$z_{score} = \frac{x_i - \mu_2}{\sigma_2}$$

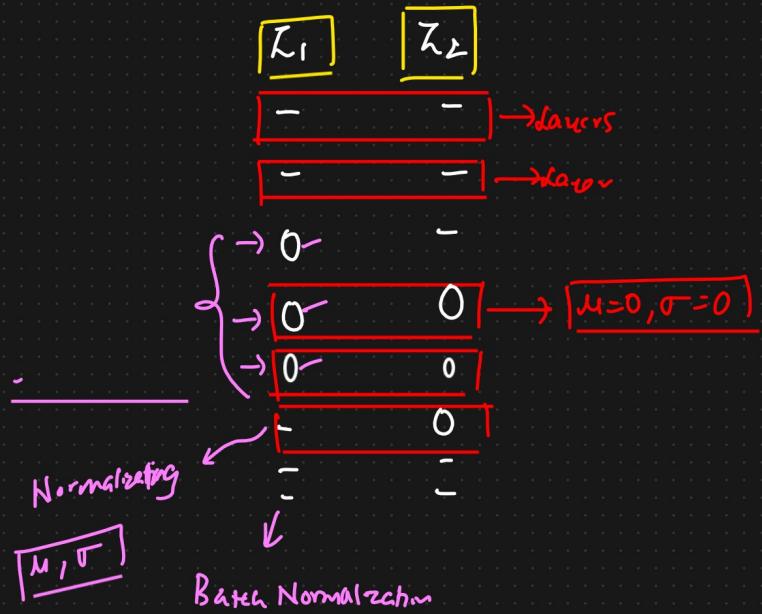
 $-$ $-$ $-$

$$z_{score} = \frac{x_i - \mu_3}{\sigma_3}$$

Layer

Normalization

$\gamma, \beta \rightarrow \text{learnable parameters}$



Normalization

\Downarrow

γ, β

$$z_1 = \gamma [w_1^T x + b_1]$$

$$y = \underline{\underline{\gamma}} \left[\frac{z_1 - \mu_1}{\sigma_1} \right] + \underline{\underline{\beta}}$$

Normalized. \Downarrow
Scale And Shift parameters

$$1) \text{ "CAT" } = [2.0, 4.0, 6.0, 8.0] \leftarrow \{ \text{Vectors} \}$$

$$2) \text{ Parameters } = \gamma = [1.0, 1.0, 1.0, 1.0] \rightarrow \text{Learned Scale} \\ \beta = [0.0, 0.0, 0.0, 0.0] \rightarrow \text{Shift} \quad \Downarrow \quad \begin{matrix} \text{Scale And} \\ \text{Shift param} \end{matrix}$$

i) Compute the mean

$$z_{\text{score}} = \frac{x_i - \mu}{\sigma}$$

$$\mu = \frac{1}{4} (2.0 + 4.0 + 6.0 + 8.0)$$

$$= \frac{20.0}{4} = 5.0$$

ii) Compute the variance (σ^2)

$$\sigma^2 = \frac{1}{4} \left[(2.0 - 5.0)^2 + (4.0 - 5.0)^2 + (6.0 - 5.0)^2 + (8.0 - 5.0)^2 \right] = 5.0$$

iii) Normalize the i/p

$$\hat{x}_i = \frac{x_i - \mu}{\sqrt{\sigma^2 + \epsilon}} \quad \epsilon = 1e^{-5} \Rightarrow \text{Avoid division by 0}$$

$$\sqrt{\sigma^2 + \epsilon} = \sqrt{5.0 + 1e^{-5}} \approx \sqrt{5.00001} = 2.236$$

$$\hat{x}_1 = \frac{2.0 - 5.0}{2.236} \approx -1.34 \quad \text{Normalized Vector}$$

$$\hat{x}_2 = \frac{4.0 - 5.0}{2.236} \approx -0.45 \quad \hat{x} = [-1.34, -0.45, 0.45, 1.34]$$

$$\hat{x}_3 = \frac{6.0 - 5.0}{2.236} \approx 0.45$$

$$\hat{x}_4 = \frac{8.0 - 5.0}{2.236} \approx 1.34$$

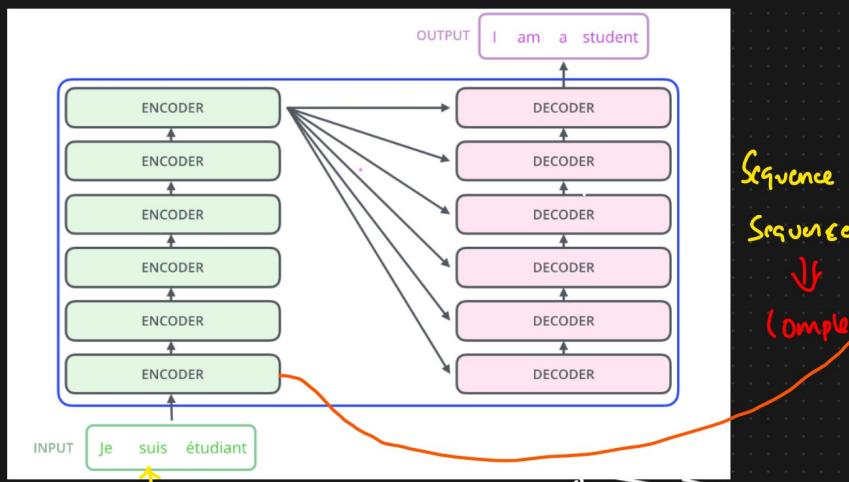
4) Scale And Shift

$$y_i = \gamma_i \hat{x}_i + \beta_i$$

$$\gamma = [1.0, 1.0, 1.0, 1.0] \quad \beta = [0.0, 0.0, 0.0, 0.0]$$

$$y = [-1.34, -0.45, 0.45, 1.34]$$

① Encoder Architecture [Research Paper]



INPUT: Je suis étudiant

Encoder: Text Embeddings + Positional Encoding \Rightarrow 512. \Rightarrow {Research paper}. Every word = 512.

Decoder: Multi Head Attention \rightarrow $z_1, z_2, z_3, z_4, z_5, z_6, z_7, z_8$ {Research paper}. $K = 64$

Residuals: $\text{Add and Norm} \rightarrow$ $\text{Feed Forward NN.} \rightarrow$ $\text{Add and Norm} \rightarrow$ $\text{Text Embeddings + Positional Encoding}$

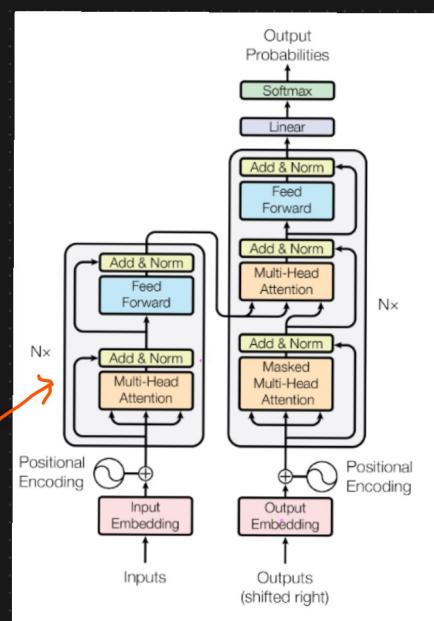


Figure 1: The Transformer - model architecture.

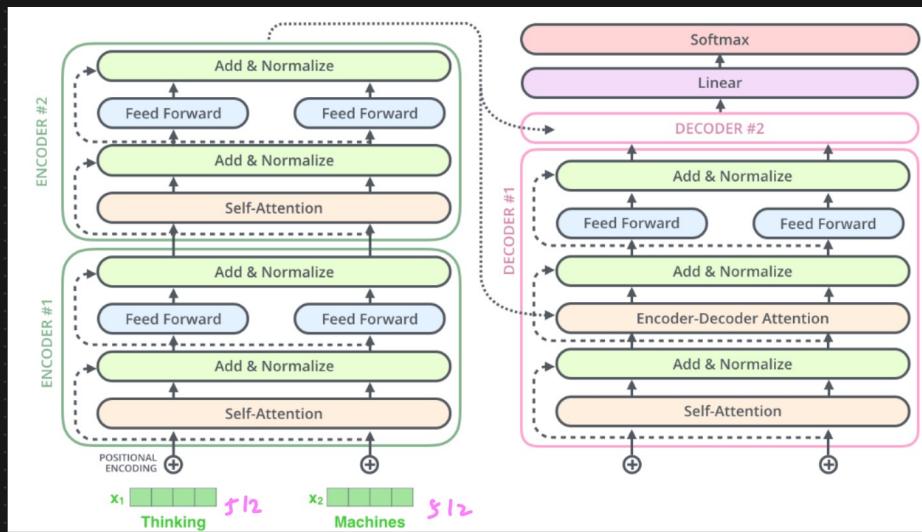
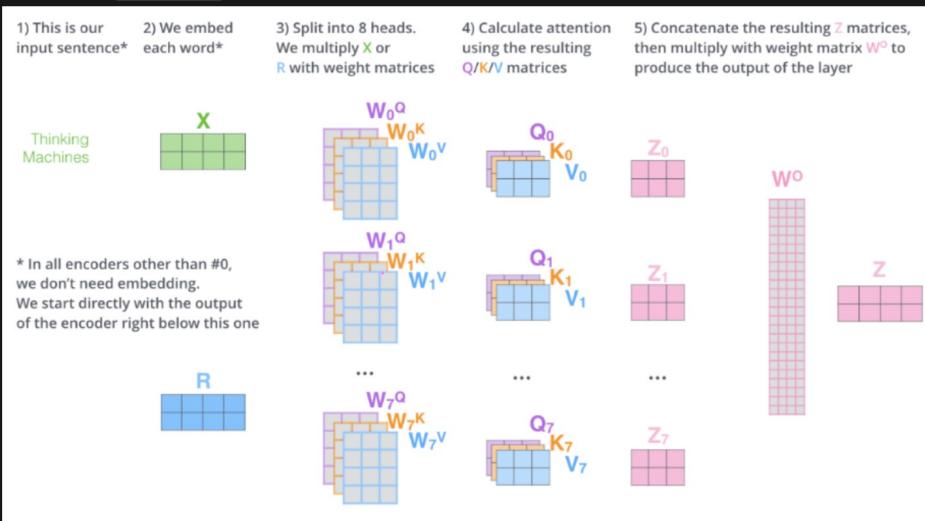
$$d_h = 64$$

$$K = 64$$

$$V = 64$$

$$\sqrt{64} = 8$$

Self Attention to Feed Forward Neural N/W



① Residual connection : Skip connection NN.

1) Addressing the Vanishing Gradient Problem

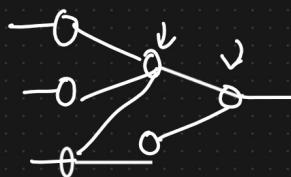
Residuals : Residual connection create a short paths for gradients to flow directly through the n/w. Gradient remains sufficiently large.

2) Improve Gradient flow

Convergence will be faster.

3) Enables Training of Deeper Networks.

① Feed Forward NN



dinner function problem
{Non linear functions}.

① Adding Non linearity

② Processing Each position Independently.

Self Attention \rightarrow (capture relationships)

FFN \rightarrow Each token representation independently.



Transforming those representation further

and allows the model to learn

Richer Representation.

ANN \Rightarrow

③ FFN \rightarrow Depth \Rightarrow Adds Depth to the Model.

Depth $\uparrow\uparrow \Rightarrow$ More learnings \rightarrow DATA

① Decoders In Transformers

3 main Components

The transformer decoder is responsible for generating the output sequence one token at a time, using the encoder's output and the previously generated tokens.

① Masked Multi Head Self Attention ✓

② Multi Head Attention (Encoder Decoder Attention)

③ Feed Forward Neural Network.

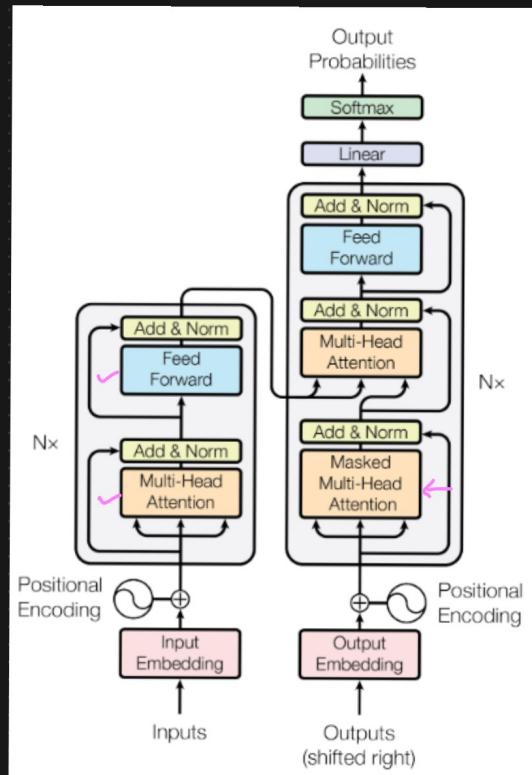
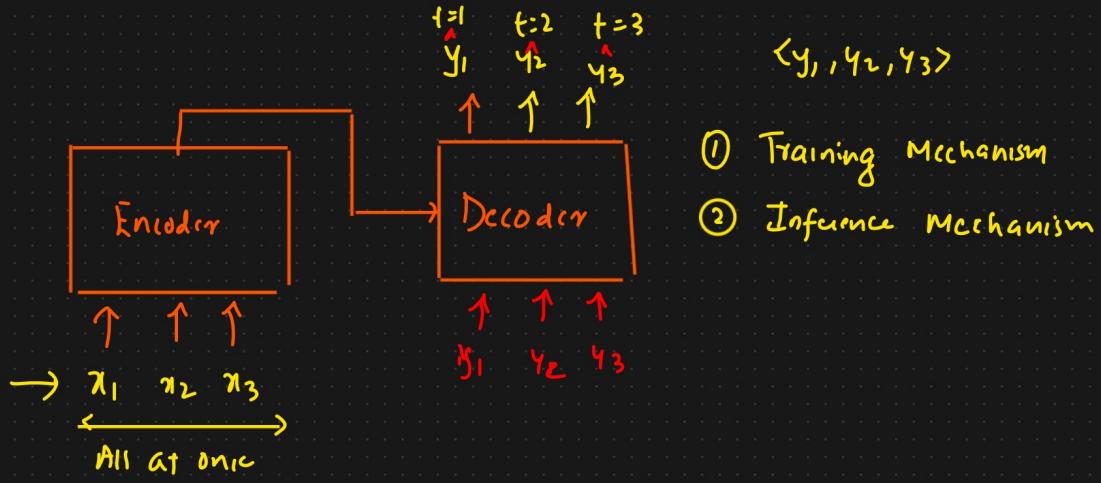


Figure 1: The Transformer - model architecture.



④ Masked Multi Head Self Attention

- ① I/P Embedding And Positional Embedding ✓ → Two padding → Sequence length equal
- ② Linear Projection for Q, K, V
- ③ Scaled Dot Product Attention
- ④ Mask Application } => Try to understand the imp. }
- ⑤ Multi Head Attention
- ⑥ Concatenation And Final Linear Projection
- ⑦ Residual Connection And Layer Normalization

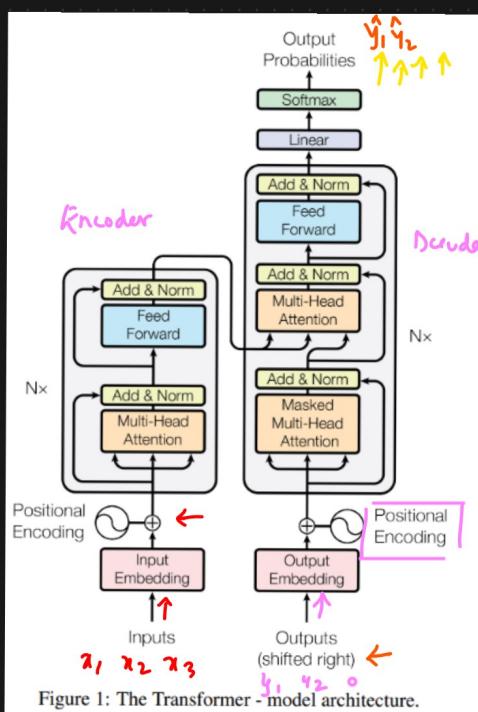
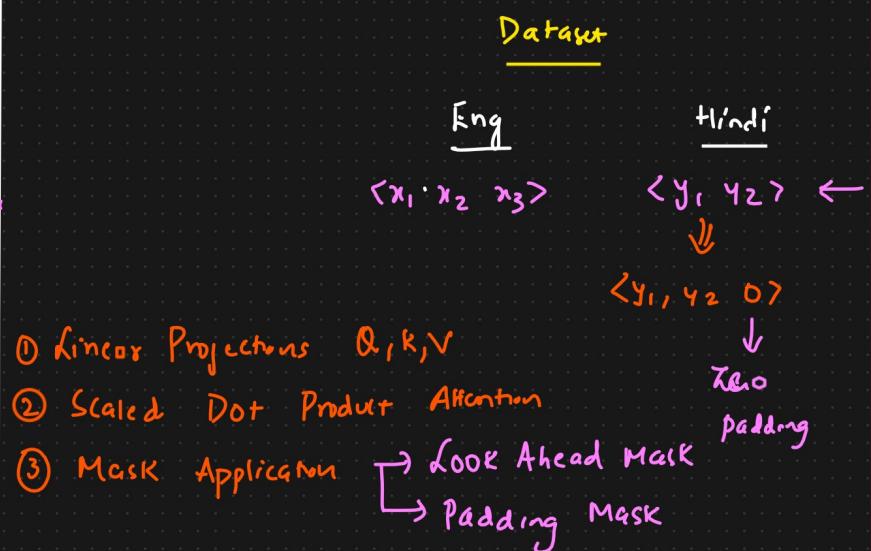


Figure 1: The Transformer -model architecture.



Masked
Multi Head
Attention.

$$[4 \ 5 \ \frac{6}{I/P} \ 7] \quad [1 \ 2 \ \frac{3}{O/P}]$$

$$[1 \ 2 \ 3 \ 0]$$

↓
4 dimension vector

i) Input Embedding and Positional Encoding

Output Embedding [STEP 1]

$$\begin{bmatrix} [0.1, 0.2, 0.3, 0.4], \\ [0.5, 0.6, 0.7, 0.8], \\ [0.9, 1.0, 1.1, 1.2], \\ [0.0, 0.0, 0.0, 0.0] \end{bmatrix}$$

$$WQ = WK = WV = I$$

Step 2 : Linear Projection for Q, K, and V.

Create query (Q), key (K) and value (V) vectors

Q = Output Embedding $\times W_Q$ = Output Embedding

K = " $\times WK$ = " "

V = " $\times WV$ = " "

$$Q = K = V = \begin{bmatrix} [0.1, 0.2, 0.3, 0.4], \\ [0.5, 0.6, 0.7, 0.8], \\ [0.9, 1.0, 1.1, 1.2], \\ [0.0, 0.0, 0.0, 0.0] \end{bmatrix}, \begin{bmatrix} 0.1 \\ 0.2 \\ 0.3 \\ 0.4 \end{bmatrix} \times K^T \begin{bmatrix} 0.5 \\ 0.6 \\ 0.7 \\ 0.8 \end{bmatrix}, \begin{bmatrix} 0.9 \\ 1.0 \\ 1.1 \\ 1.2 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

③ Scaled Dot Product Attention Calculation

$$\text{Scores} = Q \times K^T / \sqrt{dk}$$

$$= Q \times K^T / 2$$

$$\begin{bmatrix} 0.1 \times 0.1 + 0.2 \times 0.2 + 0.3 \times 0.3 + 0.4 \times 0.4, \\ 0.1 \times 0.5 + 0.2 \times 0.6 + 0.3 \times 0.7 + 0.4 \times 0.8, \\ 0.1 \times 0.9 + 0.2 \times 1.0 + 0.3 \times 1.1 + 0.4 \times 1.2 \\ 0.1 \times 0 + 0.2 \times 0 + 0.3 \times 0 + 0.4 \times 0 \end{bmatrix}$$

$$\underline{\text{Score}} : \left[\begin{bmatrix} 0.3, 0.7, 1.1, 0.0 \end{bmatrix} \quad \left[\begin{bmatrix} 0.7, 1.9, 3.1, 0.0 \end{bmatrix} \quad \left[\begin{bmatrix} 1.1, 3.1, 5.1, 0.0 \end{bmatrix} \quad \left[\begin{bmatrix} 0.0, 0.0, 0.0, 0.0 \end{bmatrix} \right] \right] \right]$$

④ Masked Application

It helps manage the structure of the sequences being processed and ensures the models behaves correctly during training and Inference

Reasons

① Handling Variable length Sequences with Padding MASK

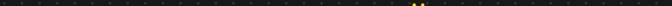
Purpose

- ① To handle sequences of different length in batch
 - ② To ensure that padding tokens, which are added to make sequences of uniform length, do not affect the model prediction

Eg: $1/1 \leftarrow \text{Sequence 1} [1, 2, 3]$

Op

o/p \leftarrow Sequence 2 [4, 5, 10] 0 is the padding token

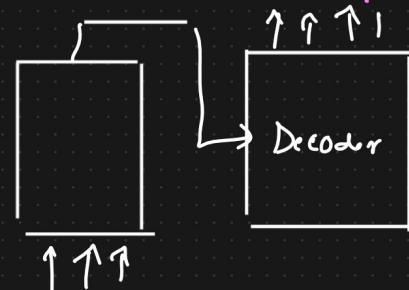
→ 100. 

lead to Incorrect or biased predictions.

A padding mask \Rightarrow The tokens are ignored.

MASKING \rightarrow Padding Mask \rightarrow Look Ahead Mask. } . Padding Mask $\begin{bmatrix} 1 & 1 & 1 \end{bmatrix}$ $\begin{bmatrix} 1 & 1 & 0 \end{bmatrix}$

② Look Ahead Mask \rightarrow Maintain Auto Regressive Property



How Anyo

① To ensure that each position in the decoder

Output sequence can only attend to

the previous position, but no future position

② Sequence \rightarrow Language Modelling, Translation

Eg: $[4, 5, 0] \rightarrow [1, 1, 0]$ \rightarrow 1D Mask.

Token 1 attends 1, 2
Attention \leftarrow to token 1, 2
Mechanism 1, 2

$$\begin{bmatrix} 1 & 1 & 0 \\ 1 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

Convert 1D to 2D Mask

For each token in the sequence,
the mask should indicate
which tokens it can attend
to.

③ Look Ahead Mask \rightarrow Decoder Output

$$\begin{bmatrix} [1, 0, 0] \\ [1, 1, 0] \\ [1, 1, 1] \end{bmatrix}$$

④ Combine Padding And Lookahead Mask

Element wise multiplication of 2 mask

Combine Mask = $\begin{bmatrix} [1, 0, 0] \\ [1, 1, 0] \\ [0, 0, 0] \end{bmatrix}$

④ Mask

$$\text{Scores} : \begin{bmatrix} [0.3, 0.7, 1.1, 0.0] \\ [0.7, 1.9, 3.1, 0.0] \\ [1.1, 3.1, 5.1, 0.0] \\ [0.0, 0.0, 0.0, 0.0] \end{bmatrix}$$

Look Ahead Mask

$$\begin{bmatrix} [1, 0, 0, 0] \\ [1, 1, 0, 0] \\ [1, 1, 1, 0] \\ [1, 1, 1, 1] \end{bmatrix}$$

Padding Masking [extended to 2D Format]

$$\begin{bmatrix} [1, 1, 1, 0] \\ [1, 1, 1, 0] \\ [1, 1, 1, 0] \\ [0, 0, 0, 0] \end{bmatrix}$$

Combined Mask: Look Ahead Mask + Padding Mask

$$\begin{bmatrix} [1*1, 0*1, 0*1, 0*0] \\ [1*1, 1*1, 0*1, 0*0] \\ [1*1, 1*1, 1*1, 0*0] \\ [1*1, 1*1, 1*1, 1*0] \end{bmatrix} = \begin{bmatrix} [1, 0, 0, 0] \\ [1, 1, 0, 0] \\ [1, 1, 1, 0] \\ [1, 1, 1, 0] \end{bmatrix} \Rightarrow \begin{bmatrix} [1, -\infty, -\infty, -\infty] \\ [1, 1, -\infty, -\infty] \\ [1, 1, 1, -\infty] \\ [1, 1, 1, -\infty] \end{bmatrix}$$

Masked Score

$$\begin{bmatrix} [0.3, -\infty, -\infty, -\infty] \\ [0.7, 1.9, -\infty, -\infty] \\ [1.1, 3.1, 5.1, -\infty] \\ [0.0, 0.0, 0.0, -\infty] \end{bmatrix}$$

Zero out the influence when the softmax is applied.

Attention weight.



Softmax

Softmax Score = Softmax (Masked Scores)

$$= \left[\begin{bmatrix} 1.0, 0.0, 0.0, 0.0 \end{bmatrix}, \begin{bmatrix} 0.3, 0.7, 0.0, 0.0 \end{bmatrix}, \begin{bmatrix} 0.1, 0.3, 0.6, 0.0 \end{bmatrix}, \begin{bmatrix} 1.0, 0.0, 0.0, 0.0 \end{bmatrix} \right]$$

Weight Sum of Value

Attention O/p = Softmax Scores * V.

Masking

Masking in the transformer architecture is essential for several reasons. It helps manage the structure of the sequences being processed and ensures the model behaves correctly during training and inference. Here are the key reasons for using masking:

1. Handling Variable-Length Sequences with Padding Mask

Purpose

To handle sequences of different lengths in a batch.

To ensure that padding tokens, which are added to make sequences of uniform length, do not affect the model's predictions.

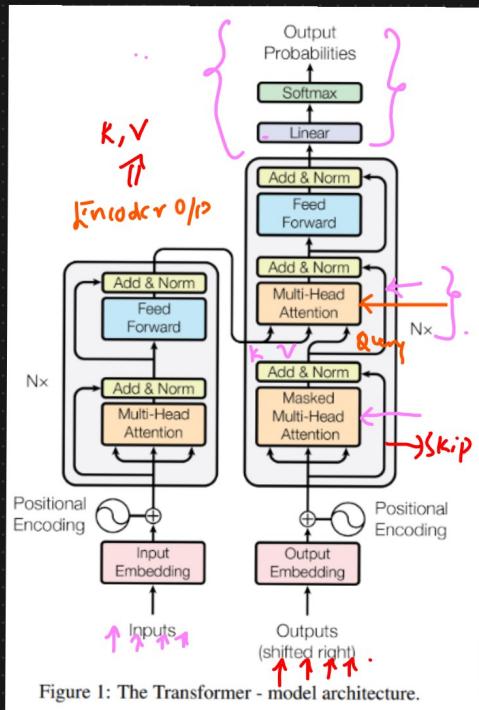
2. Maintaining Autoregressive Property with Look-Ahead Mask

Purpose

To ensure that each position in the decoder output sequence can only attend to previous positions and itself, but not future positions.

This is crucial for sequence generation tasks like language modeling and translation, where the model should not have access to future tokens when predicting the current token.

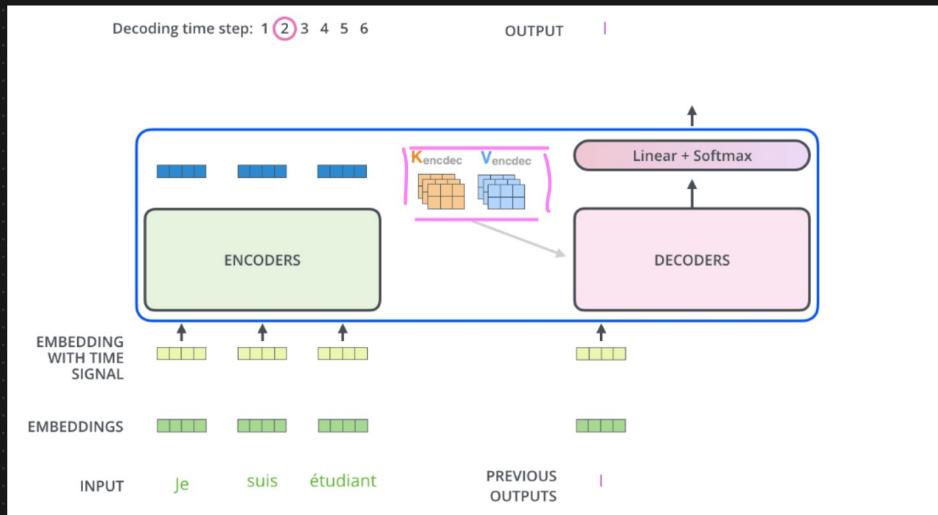
④ Encoder Decoder Multi Head Attention



- ① Encoder Out \rightarrow Set of Attention Vector K & V
- ② Masked Multihed \rightarrow Attention Vector Q {Query Vector}

These are to be used by each decoder in its
"Encoder-decoder" attention layer
↓
mention Helps the Decoder to focus on
appropriate places in the i/p Sequence

Data ref



④ The Final Linear And Softmax Layer { Vectors → o/p Word }

{ BLOG } \Rightarrow Transformers

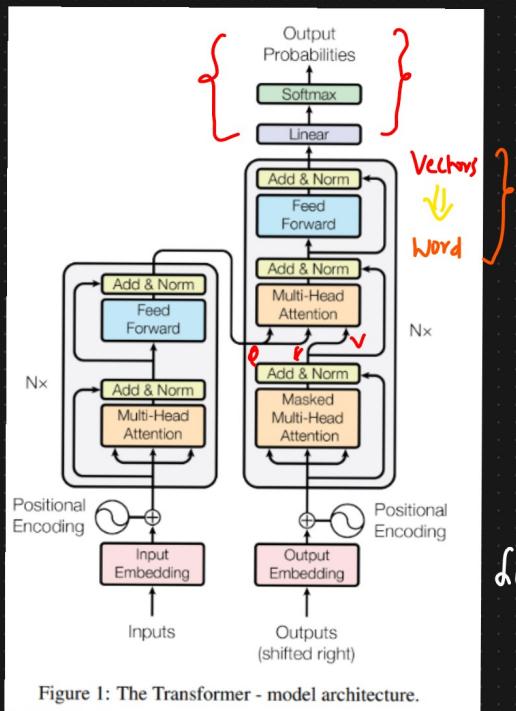
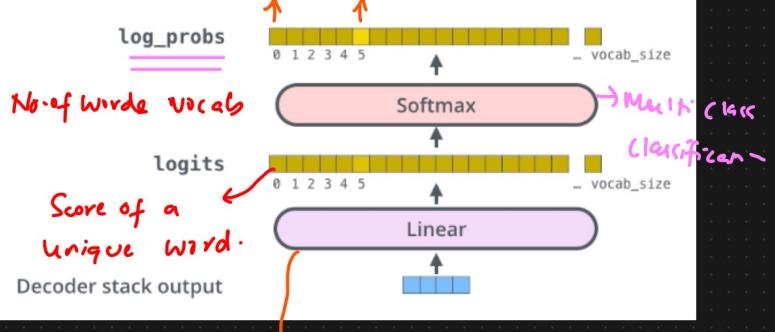


Figure 1: The Transformer - model architecture.

Which word in our vocabulary is associated with this index?

Get the index of the cell with the highest value (argmax)



linear \Rightarrow The linear layer is a simple fully connected neural net that projects the vector produced by the stack of Decoder \Rightarrow logits vector \Rightarrow

Model = 10,000 \Rightarrow Vocabulary \Rightarrow logits vector = 10000 cells wide

⑤ Softmax layer turns those Scores into probabilities (all add upto 1.0).

The cell with the highest probability is chosen, and the word associated with it is produced as the o/p. = time stamp.

Recap of Training

Output Vocabulary						
WORD	a	am	I	thanks	student	<eos>
INDEX	0	1	2	3	4	5

Output Vocabulary						
WORD	a	am	I	thanks	student	<eos>
INDEX	0	1	2	3	4	5
One-hot encoding of the word "am"						
0.0	1.0	0.0	0.0	0.0	0.0	0.0

Merci → Thanks

I am a student <eos>.

